

Two-Step Quantization for Low-bit Neural Networks

2018-10-15

Overview

- CVPR 2018
- Peisong Wang, Qinghao Hu, Yifan Zhang, Chunjie Zhang, Yang Liu, and Jian Cheng
- Two-Step量化：
 - 1 使用全精度weight量化activation，类似HWGQ但做了稀疏化
 - 2 抛弃之前全精度weight，量化transformation function (Low-bit)
 - 3 两步中间做了OTWA初始化
 - 4 使用Asymmetric Transformation Function Learning

Two-Step Quantization

$$\begin{aligned} & \underset{\{W_l\}}{\text{minimize}} && \mathcal{L}(Z_L, y) \\ & \text{subject to} && \hat{W}_l = Q_w(W_l) && \text{step2} \\ & && Z_l = \hat{W}_l \hat{A}_{l-1} && (2) \\ & && A_l = \psi(Z_l) && \text{step1} \\ & && \hat{A}_l = Q_a(A_l), \text{ for } l = 1, 2, \dots, L \end{aligned}$$

For the first step, all weights are full-precision values and we use the proposed **sparse quantization method to quantize all activations** into low-bit format. After the first step, only the learned codes A_l are kept while **the learned weights are discarded**, hence the name of **code learning**.

For the second step, we will learn the **transformation function** from A_{l-1} to A_l , with low-bit constraints.

Two-Step Quantization

3.1. Sparse Quantization for Code Learning

Two-Step: 对重要的参数求解，其余参数稀疏化为0，量化函数变为：

优化方程变为：

$$Q_{\epsilon}^*(x) = \underset{Q_{\epsilon}}{\operatorname{argmin}} E_{x \sim \mathcal{N}(0,1), x > \epsilon} [(Q_{\epsilon}(x) - x)^2] \quad (6)$$

参数的稀疏化在batch normalization之后，使用给定的 θ 让分布满足：

$$\Phi(\epsilon) = P(x \leq \epsilon) = \theta \quad (7)$$

where $\Phi(x)$ is the cumulative distribution function of standard normal distribution.

Two-Step Quantization

3.2. The Transformation Function Learning

$$\begin{aligned} & \underset{\Lambda, \hat{W}}{\text{minimize}} \quad \|Y - Q_{\epsilon}(\Lambda \hat{W} X)\|_F^2 \\ & = \underset{\{\alpha_i\}, \{\hat{w}_i^T\}}{\text{minimize}} \quad \sum_i \|y_i^T - Q_{\epsilon}(\alpha_i \hat{w}_i^T X)\|_2^2 \end{aligned} \quad (8)$$

non-linear least square regression problem

分解成多个子问题

$$\underset{\alpha, \hat{w}}{\text{minimize}} \quad \|y - Q_{\epsilon}(\alpha X^T \hat{w})\|_2^2 \quad (9)$$

为求解子问题引入辅助变量 z 和惩罚系数 λ （常量，趋近无穷时结果收敛于子问题的结果）将原函数松弛

$$\underset{\alpha, \hat{w}, z}{\text{minimize}} \quad \|y - Q_{\epsilon}(z)\|_2^2 + \lambda \|z - \alpha X^T \hat{w}\|_2^2 \quad (10)$$

Two-Step Quantization

Solving α and \hat{w} with z fixed.

Solving α and \hat{w} with z fixed. When z is fixed, the optimization problem of (10) becomes to

$$\underset{\alpha, \hat{w}}{\text{minimize}} \quad J(\alpha, \hat{w}) = \|z - \alpha X^T \hat{w}\|_2^2 \quad (11)$$

By expanding Eq. (11), we have

$$J(\alpha, \hat{w}) = z^T z - 2\alpha z^T X^T \hat{w} + \alpha^2 \hat{w}^T X X^T \hat{w} \quad (12)$$

By setting $\partial J / \partial \alpha = 0$, the optimal value of α is given by

$$\alpha^* = \frac{z^T X^T \hat{w}}{\hat{w}^T X X^T \hat{w}} \quad (13)$$

Substituting α^* in equation (12), we can get

$$\hat{w}^* = \underset{\hat{w}}{\text{argmax}} \frac{(z^T X^T \hat{w})^2}{\hat{w}^T X X^T \hat{w}} \quad (14)$$

将z视为常量，则10变为11

11展开变成一二次函数

求J对于a的极小值点（求导）

带入 α^* ，得到 w^* 并使得右侧函数最大化，在这里n-bit量化需要遍历 2^{mn} 个可行解，文章用了近似求解代替，只需要查找 $m2^n$ 个点（在这里n很小）

Two-Step Quantization

Solving z with α and \hat{w} fixed.

$$\underset{z_i}{\text{minimize}} \quad (y_i - Q_\epsilon(z_i))^2 + \lambda(z_i - v_i)^2 \quad (15)$$

w, a 视为常量（上一轮已解出）， **one-dimensional problems**
， 10变为15

松弛 Q_x :

$$\tilde{Q}_\epsilon(x) = \begin{cases} M & x > M, \\ x & 0 < x \leq M, \\ 0 & x \leq 0 \end{cases} \quad (16)$$

$$z_i^{(0)} = \min(0, v_i) \quad (17)$$

$$z_i^{(1)} = \min(M, \max(0, \frac{\lambda v_i + y_i}{1 + \lambda})) \quad (18)$$

$$z_i^{(2)} = \max(M, v_i) \quad (19)$$

按照 Q_x 分三段讨论得出 z

Two-Step Quantization

Initialization of α and \hat{w} using Optimal Ternary Weights Approximation (OTWA).

在code learning之后用0 1 -1对weight和alpha进行2-bit初始化

$$\begin{aligned} & \underset{\alpha, \hat{w}}{\text{minimize}} \quad \|w - \alpha \hat{w}\|_2^2 \\ & \text{subject to} \quad \alpha > 0 \\ & \quad \quad \hat{w} \in \{-1, 0, +1\}^m. \end{aligned} \quad (20)$$

所以在里保留绝对值最大的top r个weight进行sign得到w

$$\hat{w}_j = \begin{cases} \text{sign}(w_j) & \text{abs}(w_j) \text{ in top } r \text{ of } \text{abs}(w) \\ 0 & \text{others} \end{cases} \quad (22)$$

$$\begin{aligned} \alpha^* &= \frac{w^T \hat{w}}{\hat{w}^T \hat{w}} \\ \hat{w}^* &= \underset{\hat{w}}{\text{argmax}} \frac{(w^T \hat{w})^2}{\hat{w}^T \hat{w}} \end{aligned} \quad (21)$$

在这里可以确定a, w在22确定（这里是让w尽可能少但尽可能大）

Two-Step Quantization

Initialization of α and \hat{w} using Optimal Ternary Weights Approximation (OTWA).

Algorithm 1 OTWA weight approximation

Input: weight matrix $W \in R^{k \times m}$

Output: $\hat{W} \in \{+1, 0, -1\}^{k \times m}$

Output: k floating point scaling factors $\{\alpha_i\}_{i=1}^k$

```
1:  $J(r) \leftarrow 0$ , for  $r = 1, \dots, m$ 
2: for  $i = 1, \dots, k$  do
3:    $v \leftarrow \text{abs}(w_i)$ 
4:   sort  $v$  in decreasing order
5:   for  $r = 1, \dots, m$  do
6:      $s \leftarrow \sum_{j=1}^r v(j)$ 
7:      $J(r) \leftarrow s^2/r$ 
8:   end for
9:    $r^* \leftarrow \text{argmax}_r J(r)$ 
10:  get  $\hat{w}_i^*$  according to equation 22
11:  get  $\alpha_i^*$  according to equation 21
12: end for
```

Asymmetric Transformation Function Learning.

$$\underset{\Lambda, \hat{W}}{\text{minimize}} \quad \|Y - Q_\epsilon(\Lambda \hat{W} \tilde{X})\|_F^2 \quad (23)$$

为了消除因为各层之间独立量化造成的误差的累积，文章将transformation function learning里面的X换成了quantized network中上一层的activation

Two-Step Quantization

4. Experiments On ImageNet

Table 1. The optimal thresholds and step values given different sparse ratios for 2-bit sparse quantization.

θ	50%	56.25%	62.5%	68.75%	75%
ϵ	0.00	0.16	0.32	0.49	0.68
Δ	0.5388	0.5914	0.6487	0.7139	0.7889

56.25%-75%

Table 2. Two-bit activation quantization comparison. Our sparse quantization method, denoted by SQ, is conducted under different sparsity.

Model	Sparsity (%)	Top-1 (%)	Top-5 (%)
AlexNet	50.00	58.5	81.5
HWGQ [2]	50.00	55.8	78.7
SQ-1	56.25	58.2	80.7
SQ-2	62.50	59.0	81.3
SQ-3	68.75	58.9	80.8
SQ-4	75.00	57.9	79.8

4.2. Sparse Quantization Results

Table 3. The accuracy of our two-bit activation and ternary weight quantization models before and after fine-tuning. Results are reported under different activation sparsity.

Model	Before Fine-tune		After Fine-tune	
	Top-1	Top-5	Top-1	Top-5
TFL-SQ-1	52.7	76.6	-	-
TFL-SQ-2	55.1	78.4	58.0	80.5
TFL-SQ-3	54.7	78.1	-	-
TFL-SQ-4	54.3	77.4	56.7	79.0

Two-Step Quantization

4.3. Transformation Function Learning Results




Figure 1. Top-1 accuracy after each layer is quantized for different transformation function learning settings. The results are conducted based on the SQ-2 AlexNet mode without fine-tuning.

- **OTWA**: Weight ternarization using OTWA;
- **TFL-rand**: Asymmetric transformation function learning initialized by random ternary variables;
- **TFL-sym**: Symmetric transformation function learning, initialized by OTWA;
- **TFL-asym**: Asymmetric transformation function learning, initialized by OTWA;

Two-Step Quantization

4.4. Comparison with the stateoftheart

Table 4. Comparison with the state-of-the-art low-bit quantization methods on AlexNet. The accuracy gap to the full-precision model is also reported.

Model	Top-1	Top-5	Top-1 gap	Top-5 gap
AlexNet[2]	58.5	81.5	0	0
XNOR[24]	44.2	69.2	-12.4	-12.3
BNN[30]	46.6	71.1	-11.9	-10.4
DOREFA[38]	47.7	-	-8.2	-
HWGQ[2]	52.7	76.3	-5.8	-5.2
TSQ (ours)	58.0	80.5	-0.5	-1.0

Table 5. Comparison between our quantized VGG-16 model and the full-precision counterparts.

Model	Top-1	Top-5	Top-1 gap	Top-5 gap
VGG-16 [29]	71.1	89.9	0	0
VGG-16-BN	69.6	89.6	-1.5	-0.3
TSQ (ours)	69.1	89.2	-2.0	-0.7

4.5. Results on CIFAR10

Table 6. Comparison with the state-of-the-art low-bit quantization methods on CIFAR-10. The bit-width for activations and weights are given.

Activation	Weights	Method	error (%)
Full	Full	VGG-Small	6.82
Full	Binary	BinaryConnect [6]	8.27
Full	Ternary	TWN [23]	7.44
Binary	Binary	BNN [14]	10.15
2-bit	Binary	HWGQ [2]	7.49
2-bit	Ternary	TSQ (ours)	6.51